

Chapter 1

Introduction

Da-Wei Chang

CSIE.NCKU

Objectives

- Introduce the major operating system components and functions
- To provide coverage of basic computer system organization
- Many of the topics will be discussed in detail in the following chapters

Outline

- What Is an Operating System
- Computer-System Organization/Architecture
- Operating-System Structure/Operations
- Process Management
- Memory Management
- Storage Management
- Protection and Security
- Virtualization
- Basic Kernel Data Structures
- Computing Environments
- Open-Source Operating Systems


What is an Operating System?

- A **program** that acts as an intermediary 中介 between **user programs/applications** and the **computer hardware**
- Operating system goals
 - Execute user programs
 - Make the computer system convenient to use
 - Applications do not have to deal with HW directly
 - OS ensures **efficient** resource **sharing**

Computer System Structure

- A computer system can be divided into 4 components (from bottom to top)
 - **Hardware** – provides basic computing resources
 - CPU, memory, I/O devices
 - **Operating system**
 - Controls and coordinates use of hardware among various applications and users
 - **System and application programs** – solve the computing problems of the users
 - Word processors, compilers, web browsers, database systems, video games
 - **Users**
 - People, machines, other computers

Four Components of a Computer System



Operating System Definition

- **What is an operating system?**
- No universally accepted definitions
- “Everything a vendor ships when you order an operating system” is good approximation
 - But varies widely
- “The one program running at all times on the computer” (i.e., the **kernel**) Everything else is either a **system program** (ships with the operating system) or an **application program**

Operating System Definition


- Two points of view (in the text book)
 - OS is a **resource allocator**
 - Manages all resources
 - Applications do not have to manage resources directly
 - OS hides HW interface
 - Decides between conflicting requests for **efficient** and **fair** resource use
 - OS is a **control program**
 - Controls execution of programs to prevent **errors** and **improper use** of the computer
- Other definitions
 - A program that can monitor the other programs
 - A program that allows the other programs to **share** resources in a **controlled** manner
 - CPU, memory, IO devices....

Computer Startup

- Who loads the OS ? Bootstrap program 負責加載OS
- Bootstrap program is loaded at power-up or reboot
 - Typically stored in ROM or EEPROM, generally known as firmware 韌體
 - Jobs
 - Initializes all aspects of system
 - Loads operating system kernel and starts OS execution
- A bootstrap program is also called a boot loader
- Who loads the boot loader?
 - In PC, BIOS loads the (in-disk) boot loader
 - In many other platforms, boot loader is placed at a predefined memory address (in ROM or EEPROM)
 - Hardware jumps to the boot loader directly

Computer System Organization

- Computer-system operation
 - One or more CPUs, device controllers connect through common bus providing access to shared memory 共用同一條bus
 - Concurrent execution of CPUs and devices competing for bus cycles 同時進行的工作會競爭bus



Computer-System Operation

- I/O devices and the CPU can execute concurrently.
- Each device controller is in charge of a particular device type.
- Each device controller has a local buffer.
- Writing data to device *a part of OS*
 - (Device driver running on the) CPU moves data from main memory to local buffers 設備驅動
 - Device controller writes the data from the local buffer to the device
 - Device controller informs CPU that it has finished its operation by generating an *interrupt*
- One of the jobs of an operating system is to manage interrupts.

Common Functions of Interrupts

中斷CPU正在做的事


中斷後，控制權交給ISR

- Interrupt transfers control to the **interrupt service routine (ISR)**, generally through the ***interrupt vector (table)*** 一個表格，每個中斷對應到一個地址
 - *interrupt vector* contains the addresses or instructions of all the service routines
 - X86: addresses
 - ARM: instructions
- The processor must save the address of the **interrupted** instruction before jumping to the ISR 被中斷的指令要存
- Generally, interrupts are raised by IO device controllers
 - **Hardware** interrupts

Common Functions of Interrupts

- However, **software** can also generate/raise interrupts
 - Software interrupts (or **traps**)
- A *trap* is a software-generated interrupt caused either by an **error** or a **user request**
 - Error
 - Division by zero, invalid memory access...
 - User request
 - Requests for OS services
 - **syscall/int** instruction in x86
 - **svc/swi** instruction in ARM
- An operating system is **interrupt driven!**


OS靠中斷驅動




Interrupt Handling

- 1 • Save the CPU state (*registers*, including program counter)
 - Done by HW or OS
- 2 • Locate the corresponding Interrupt Service Routine (ISR) or interrupt handler
 - How do you know the address of the ISR?
 - **Interrupt vector table** (IVT)
- 3 • Execute the ISR
- 4 • Continue the interrupted work by restoring the CPU state


Saving/Restoring CPU State – An Example




Interrupt-drive I/O Cycle



Interrupt Timeline



Two I/O Methods



Direct Memory Access (DMA)

- Who moves the data between memory and device buffer?

- CPU? device controller?


 DMA Direct Memory Access

- Device controller transfers blocks of data from local buffer directly to main memory **without CPU intervention**
 - Raise interrupt when DMA transfer is done

DMA控制器不經由
CPU，把資料搬到
memory

Storage Structure

- Storage systems organized in **hierarchy**
 - The layers differ in **Sizes, Speed, Cost, Volatility...**



Storage Structure

- Primary Storage
 - Main memory –the only large storage media that the CPU can access **directly** (e.g., via load/store instructions)
- Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity
 - Magnetic disks (HDD)
 - metal or glass platters covered with magnetic recording material
 - disk surface is logically divided into **tracks/sectors**
 - the most common secondary-storage device
 - Nonvolatile Memory (NVM)
 - SSD (Solid State Drive)
 - Becoming more popular as capacity increases and price drops
 - *Various technologies: **flash memory**, PCM, RRAM, STT-MRAM*
- Tertiary storage
 - Usually for storing backup copies

Caching

- *Caching* – information in use copied from slower to faster storage temporarily
 - main memory can be viewed as a *cache* for secondary storage
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache is smaller than the storage being cached
 - Cache management is an important design problem
 - Cache size and replacement policy
- Performed at **many levels** in a computer system (in hardware, operating system, software)

Performance of Various Levels of Storage

Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	DRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	25,000 - 50,000	5,000,000
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

Movement between levels of storage hierarchy can be **explicit** or **implicit (to software)**

隱性


- SW-managed caching (explicit) register, memory, disk
- HW-managed caching (implicit). cache

Hardware

Software

顯性


Migration of Data A from Disk to Register



Computer-System Architecture


- Many systems use general-purpose processors (from PCs to mainframes)
 - Some systems have special-purpose processors as well
- Multi-processor systems growing in use and importance
 - Also known as parallel systems, tightly-coupled systems
 - Advantages
 1. Increased throughput
 2. Increased reliability – graceful degradation or fault tolerance
 3. Economy of scale
 - Cheaper than multiple single-processor systems due to sharing of resources (e.g., memory, secondary storage, IO devices...)
 - Types
 1. Asymmetric Multiprocessing
 - Not all the processors are treated equally
 2. Symmetric Multiprocessing

Symmetric Multiprocessing Architecture




Multicore Systems

- **Multicore**: multiple **compute cores** on a single chip
 - More efficient than multiple single-core chips
 - On-chip communication is faster than cross-chip communication



A dual-core design with two cores on the same chip

Graceful Degradation




suppose a four-core system

Non-Uniform Memory Access (NUMA)

- Adding too many CPUs on multiprocessor systems will cause
 - contention for **system bus**
 - Since each CPU needs to access main memory
 - Performance begins to degrade
- Alternative approach: **NUMA**
 - Provide each CPU with its own local memory
 - Accessed via a local bus (next slide)
 - Increasingly popular on servers and high-end computing systems

Non-Uniform Memory Access (NUMA)



- Problem: memory *access time* is not the same
 - remote mem. access latency > local mem. access latency
 - OS need to manage tasks/data carefully
 - Co-locates a task and its data

Clustered Systems


聚集的

- A *clustered system*
 - Multiple computers connected via a *local area network (LAN)* or a *faster interconnection network*.
- The key of a clustered system is *high availability*
 - *Fault tolerant*: suffer a failure of any single component and continue operation
 - *Graceful degradation*: continue providing service *proportional* to the level of surviving hardware

General Structure of a Clustered System

An example

*HP C7000 G3 BladeSystem
-15 nodes, 320 cores, 10 Gb Ethernet



<https://ctoservers.co.uk/hp-c7000-g3-baladesystem-w15-x-bl460c-g8--320-cpu-cores1tb-ram-10gbe---cad-cfd-lsdyna--ansys-2101-p.asp>


Clustered Systems

- ***Asymmetric Clustering*** 待命
 - Some machines are in ***hot-standby mode*** while the others are running applications. 只有部分在工作
 - A hot-standby machine (i.e., does nothing but) monitors the other machines and becomes active if one server fails.
- ***Symmetric Clustering***
 - Machines run applications while monitoring each other
 - more efficient as it uses all available hardware

Multiprogramming & Timesharing

- **Multiprogramming** is needed for **efficiency**
 - a single job cannot keep CPU and I/O devices busy at all times
 - E.g., CPU idle if the job is waiting for IO CPU在等I/O是可以先做其他工作
 - Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - increase CPU utilization 一部分jobs放到
 - A subset of total jobs in system is kept in memory memory
 - When one in-memory job has to wait (for I/O for example), OS switches to another job in the memory
 - Other jobs are kept in the **on-disk job pool** 剩下的放在job pool
 - A job is selected and put in memory via **job scheduling**

Memory Layout for a Multiprogrammed System



Multiprogramming & Timesharing

- **Timesharing (multitasking)** is logical **extension** to multiprogramming in which CPU switches jobs so **frequently** that users can interact with each job while it is running, creating **interactive** computing
 - **Response time** should be < 1 second
 - If several jobs ready to run at the same time \Rightarrow **CPU scheduling**
 - differences with job scheduling? Job scheduling is 挑選job載入到 RAM , CPU scheduling 是從RAM
 - If processes don't fit in memory, swapping moves them in and out to run
 - Virtual memory allows execution of ready queue 挑選process來執行 memory
 - Allows a program that is larger than the physical memory to run

Operating System Operations

- An OS must prevent user jobs/processes from damaging the normal operations of the OS
- **Dual-mode operation** allows OS to protect itself
 - **Processor** provides **user mode** and **kernel mode**
 - User processes run in user mode (non-privileged mode)
 - OS runs in kernel mode (privileged mode)
 - Some instructions designated as **privileged**, can only be executed in kernel mode
 - E.g. I/O control, timer, special registers, interrupt related instructions...
 - a **mode bit** is provided
 - to distinguish when system is running user code or kernel code
 - A **system call** changes the mode to the **kernel mode**
 - return from the system call resets the mode to **user mode**


Transition from User to Kernel Mode

At system boot time, the HW starts in the kernel mode.

- after system initialization, the OS run applications in user mode.


When a user process needs an OS service, it issues a [system call](#)
e.g., read a file, send a packet...

System call flow---



Transition from User to Kernel Mode

- Modes are provided by HW (CPU)
 - Intel's 8088 CPU has only a single mode
 - The OS (e.g., MS-DOS) can be crashed by a user program
 - Modern x86 has (more than) 4 modes/rings



Timers

其實OS跟其他東西都一樣是執行
在CPU上

- Conceptually, user processes run **on top of the OS**
 - *In fact, they run on the CPU directly!!!*
- How to ensure that the OS maintains control over the CPU?
那怎麼確定永遠是OS在控制CPU 用Timer避免無窮迴圈
- Prevent an infinite loop in a process from hogging resources
 - By setting up a timeout period for the process
 - Done by **timers**
 - Operating system decrements the time count for the process periodically (whenever a **timer interrupt** is generated)
 - When the count becomes zero (i.e., the current process has used up its time period)
 - hand the CPU to another process

Process Management

- A process is a program **in execution**. It is a unit of work within the system.
 - Program is a *passive entity*, process is an *active entity*.
- Process needs resources to accomplish its task
 - CPU, memory, I/O, files 回收
- Process termination requires reclaiming reusable resources
- **Single-threaded** process has one program counter specifying location of **next instruction** to execute 單線程
 - Process executes instructions sequentially, one at a time, until completion
- **Multi-threaded** process has multiple program counters
 - One program counter per thread 多線程
- Typically, a system has many processes running concurrently on one or more CPUs
 - Concurrency by multiplexing the CPUs among the processes/threads

Process Management Activities

- Activities related to process management
 - Creating and deleting processes
 - Suspending and resuming processes
 - Providing mechanisms for process synchronization
 - Providing mechanisms for process communication
 - Providing mechanisms for deadlock handling
- We will discuss the above topics later

Memory Management

- Data have to be put in memory before being processed
- Instructions have to be put in memory before being executed
- Memory management determines what is in memory
 - Goal: optimizing CPU utilization and computer's response to users
- Memory management activities
 - Allocating and deallocating memory space as needed
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory

Storage Management

- OS provides uniform, logical view of information storage
 - Different types of storage devices (e.g., solid-state drives, hard disk drives, tape drives...)
 - Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)
 - Abstracts physical properties to logical storage unit
 - file 不管底層儲存裝置是什麼，我們看到的都是檔案
 - OS stores files into the storage devices

Storage Management

- **File-System Management (provide file interface)**
 - File content is determined by its creator
 - Free format – e.g. text files...
 - Fixed format – e.g. executable files...
 - Files usually organized into **directories**
 - **Access control** on most systems to determine **who** can access **what**
 - OS activities include
 - Creating and deleting files and directories
 - Primitives to manipulate files and dirs
 - ^{原語}read/write/append files, set/get file status, set/get file permissions

Storage Management

- **Mass-Storage Management (focus on efficient operation)**
 - Usually, disks
 - Mass storage devices are used to store
 - data that do not fit in main memory, or
 - data that must be kept for a “long” period of time
 - OS activities
 - Free-space management (bitmap? free list?)
 - Storage allocation
 - Disk scheduling
 - Critical to system performance
 - Because disk is usually the **performance bottleneck** 硬碟通常是效能瓶頸
 - Latency of a computer operation usually depends on disk subsystem and its algorithms

I/O Subsystem

- One purpose of OS is to hide peculiarities of hardware devices from the user
- I/O subsystem
 - Drivers 驅動程式
 - Manages specific hardware devices 只有驅動程式是針對特定設備 (nvidia)
 - In an OS, only drivers are device-specific
 - Have knowledge about its own device
 - OS provides a general driver interface to allow easy cooperation of OS and its drivers
 - Memory management of I/O including
 - buffering (storing data temporarily while it is being transferred)
 - caching (storing parts of data in faster storage for performance)
- Description here is also valid for storage devices
 - Disks are IO devices

Protection and Security


- **Protection** – mechanisms for controlling **access** (of processes or users) **to resources**
 - Distinguish between authorized and unauthorized usage
- **Security** – defense of the system against internal and external attacks
 - Huge range, including denial-of-service, viruses...

Protection and Security

- Systems generally first distinguish among **users**, to determine who can do what
 - Each user has a user identifier (**user ID**)
 - User IDs are associated with all files & processes to determine access control
 - **Process owners** and **file owners** are all users...
 - Group identifier (**group ID**) allows set of users to be defined and managed, also associated with each process or file **批量管理**
 - 特權提升
– **Privilege escalation** allows user to change to effective ID with more rights *temporarily*
 - **Setuid** in UNIX (**see next slide**)
 - Causes a program to run with a **user ID of the file owner**, rather than the user's ID **用file owner的id來執行**
 - » E.g., change your own password, which requires the root privilege to update the password file.
 - Might become a vulnerability if the program is attacked...

Setuid in Linux


```
cissoll> ls -al passwd  
-r-sr-sr-x 1 root sys 27228 Aug 17 2007 passwd*  
cissoll>
```



Virtualization

- **Virtualization**
 - Abstract **the hardware of a single computer** into several different execution environments
 - Create an illusion that each environment is running on its own private computer
 - An OS can be run in each environment
 - Allows OSes to run on other OSes

Virtualization



host operating system vs. guest operating system


Virtualization

- **Emulation** –
 - Simulating computer hardware in software
 - Used when source CPU type **different** from target type
 - For Example
 - We want to run the PowerPC applications on the Intel x86
 - Every instruction of PowerPC must be **translated** to the instructions on Intel x86
 - Slow...

Virtualization (Cont.)


- **Virtualization** –
 - An OS **natively compiled for a CPU**, running on another OS **also natively compiled to that CPU**
 - Example
 - Run Linux for x86 on the Windows 10 (also for x86)
 - Linux: **guest OS**
 - Windows 10: **host OS**
 - **VMM (Virtual Machine Manager)** provides virtualization services
 - Also called **Hypervisor**
 - VMM can also run **natively** without the host OS **直接跑在硬體上**
 - Example: VMware ESX and Citrix XenServer

VMM without Host OS



(a)

Without virtualization



(b)

With virtualization

VMM directly runs on hardware (do not need host OS) 55


Kernel Data Structures

- Lists, Stacks, and Queues
- Trees
- Hash Tables
- Linux data structures defined in *include* files
 - e.g., `<linux/list.h>`, `<linux/kfifo.h>`,
`<linux/rbtree.h>`


常見的kernel結構

List


- *Singly linked list*



- *Doubly linked list*



- *Circular linked list*




Stack and Queue

- **Stack:** last in, first out (LIFO)
 - Insert an item: **push**
 - Remove an item: **pop**
- **Queue:** first in, first out (FIFO)


Tree

- **General tree**
 - A parent may have an unlimited number of children
- **Binary tree**
 - A parent may have at most two children
- **Binary search tree**
 - Two children, left \leq right
 - Search performance is $O(n)$
 - **Balanced binary search tree**
 - Search performance is $O(\lg n)$



Hashing and Bitmap

- **Hash function** can create a **hash map**



- **Bitmap** – string of n binary digits representing the status of n items

Computing Environments

- Traditional Computing
- Mobile Computing
- Client-Server Computing
- Peer-to-Peer Computing
- Cloud Computing
- Real-Time Embedded Systems

Computing Environments

- Traditional Computing
 - Office environment
 - PCs connected to a network, terminals attached to mainframe or minicomputers providing batch and timesharing functionalities
 - Now portals allowing remote systems to access internal resources
 - Home networks
 - Used to be single system, then modems
 - Now firewalled, networked
- 終端連到大型主機
或是小型電腦
提供批次處理 (batch)與分時

Mobile Computing


- Computing on **handheld devices**
 - Smartphones
 - Tablets
- Issues
 - Limited memory
 - Slower processors
 - I/O constraints: small display screens, small keyboards.....
 - Limited power
- Extra features
 - GPS, accelerometers, and gyroscope
- Leaders are **Apple iOS** and **Google Android**

Client-Server Computing

Servers: serving requests, **Clients**: issuing requests

- ▶ **Application-server** provides an interface to client to request services
- ▶ **File-server** provides interface for clients to store and retrieve files

For clients, dumb terminals are replaced by PCs or smart phones




Peer-to-Peer Computing

- P2P does not distinguish clients and servers
 - Instead all nodes are considered *peers*
 - May each act as client, server or both
 - Node must join P2P network
- Advantage
 - In client-server system, the server may become a bottleneck.
In P2P, services can be provided by several nodes
- Example
 - skype

Peer-to-Peer Computing

- Node must join P2P network
 - Registers its service(s) with central lookup service on network, or
 - Broadcast requests for services and respond to requests for services via *discovery protocol*




Cloud Computing

- Delivers computing, storage, even apps as a service across a network
- Logical extension of **virtualization** (see next slide)
 - Typically based on virtualization
 - Amazon **EC2** has thousands of servers, millions of **VMs**, PBs of storage available across the Internet, pay based on usage
- Different types
 - **Public cloud** – available via Internet to anyone willing to pay
 - **Private cloud** – run by a company for the company's own use
 - **Hybrid cloud** – includes both public and private cloud components

Logical Extension of Virtualization


WWW Server

Solaris




File Server

Win10




Database Server

Linux



**WWW + File + Database
Servers**




Cloud Computing

- Categories
 - **Software as a Service (SaaS)** –
 - one or more applications available via cloud (e.g., word processor)
 - **Platform as a Service (PaaS)** –
 - software stack ready for application use via cloud (e.g., a database server)
 - **Infrastructure as a Service (IaaS)** –
 - servers or storage available via cloud (e.g., storage available for backup use)

Cloud Computing

- Cloud computing environments composed of traditional OSes, plus VMMs, plus cloud management tools
 - Internet connectivity requires security like firewalls
 - Load balancers spread traffic across multiple applications



Real-Time Embedded Systems

- Embedded computers
 - Most prevalent form of computers
 - Cars, routers, robots, printers,...
 - Tend to have very **specific** tasks
- Different OS implementation schemes
 - No OS
 - A tiny embedded OS
 - A standard OS (e.g., Linux)

Real-Time Embedded Systems

- Embedded systems sometimes run *real-time operating systems*
 - Processing must be done within a predefined time constraint (i.e., **deadlines**)
 - Otherwise, the system will fail
 - e.g., missile systems

Open-Source Operating Systems

- Many operating systems are open source
 - e.g., Linux, FreeBSD, OpenSolaris, Illumos,...