

# Chapter 5

## CPU Scheduling

Da-Wei Chang

CSIE.NCKU

1

Source: Abraham Silberschatz, Peter B. Galvin, and Greg Gagne, "Operating System Concepts", 10th Edition, Wiley.

# Outline

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Multiple-Processor Scheduling
- Thread Scheduling
- Operating Systems Examples
- Algorithm Evaluation

# Basic Concepts

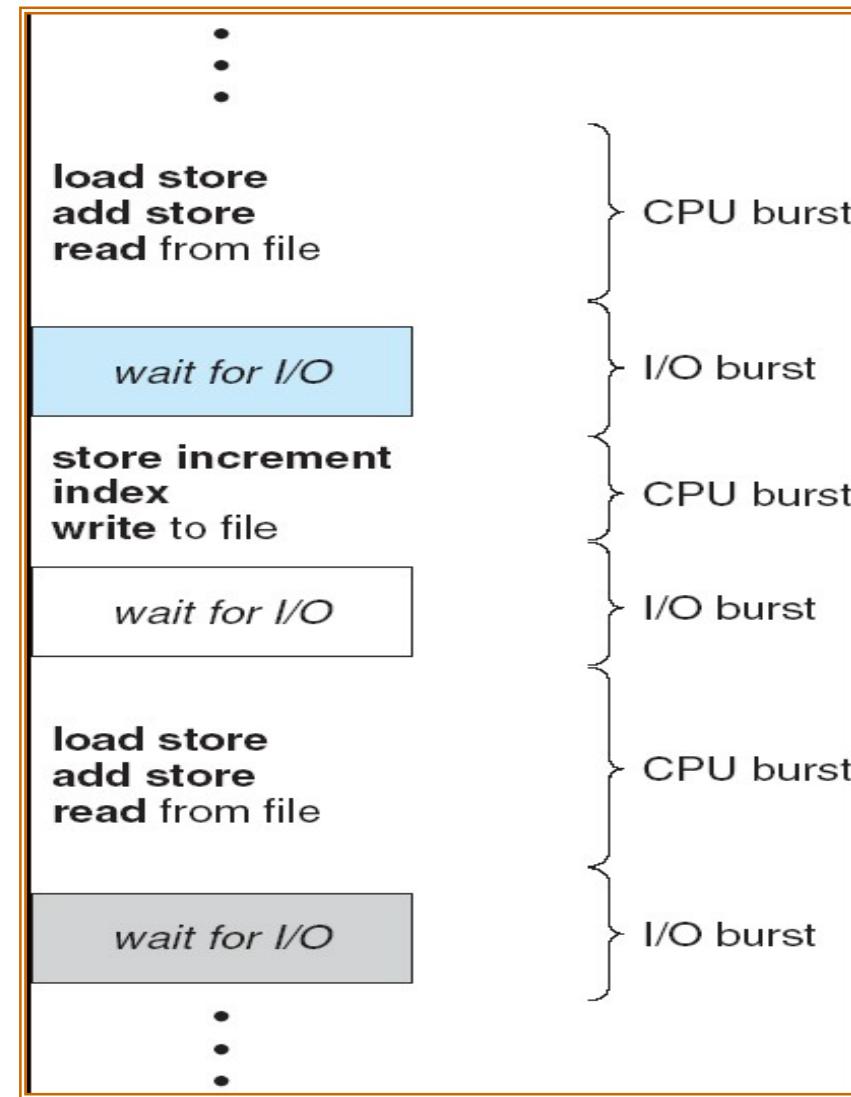
- Scheduling is a basis of multiprogramming
  - Switching the CPU among processes improves CPU utilization Schedule 是多道programming 的基礎

- CPU-I/O Burst Cycle – Process execution consists of a *cycle* of CPU execution and I/O wait
    - 突發週期
- Process的執行由CPU突發跟I/O等待週期所組成

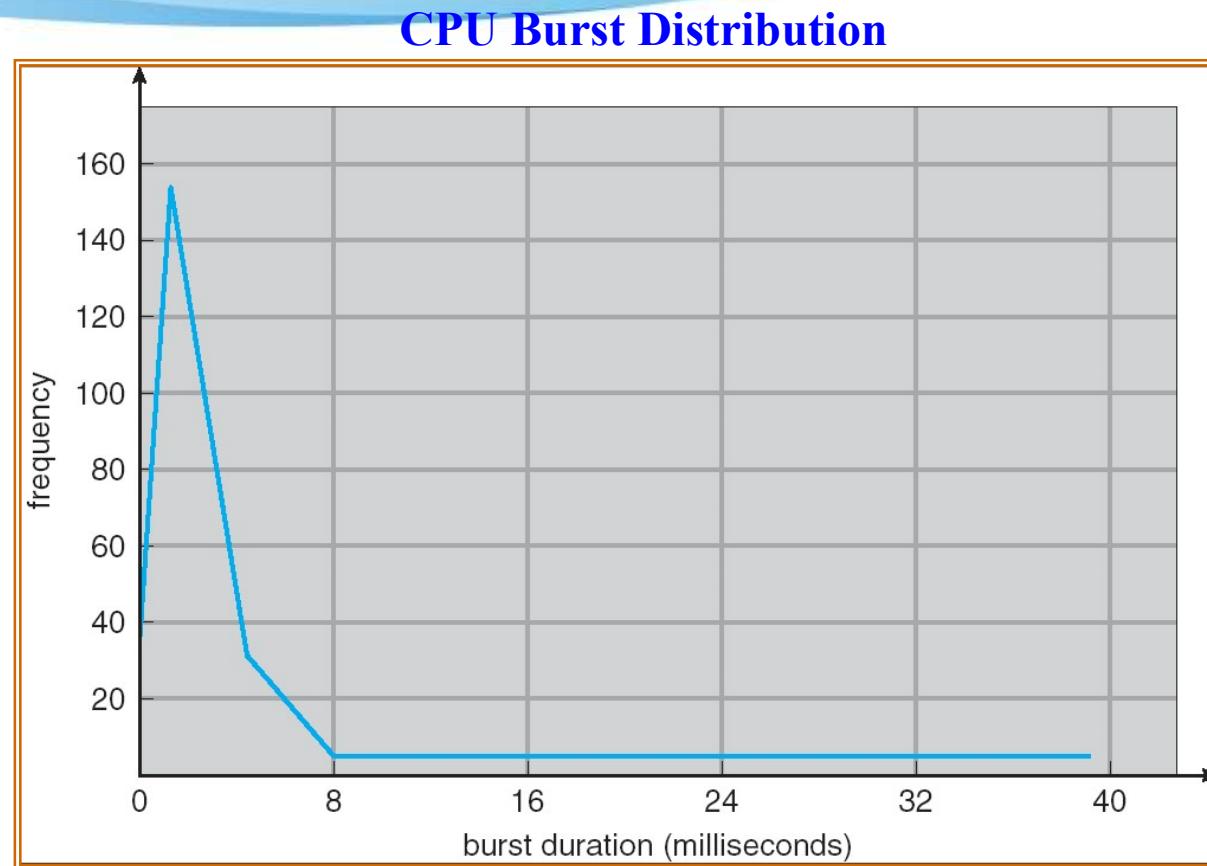
一個行程從提交到終止的整個執行過程，就是由一系列 CPU 突發和 I/O 突發交替組成的：

1. **CPU 突發 (CPU Burst)**: 行程執行計算指令。
2. **I/O 突發 (I/O Burst)**: 行程發出 I/O 請求，並進入等待狀態，CPU 被釋放。
3. **I/O 完成**: I/O 請求完成，行程從等待狀態切換回就緒狀態。
4. 下一個 **CPU 突發**: 行程再次被排程器選中，執行下一個 CPU 突發。

# Alternating Sequence of CPU and I/O Bursts



# Histogram of CPU-burst Times



A large # of short CPU bursts and a small # of long CPU bursts

兩種類型  
IO bound → many short CPU bursts, few long CPU bursts  
CPU bound → more long CPU bursts

# CPU Scheduler

- Short term scheduler
- **Selects** among the processes in memory that are ready to execute, and allocates the CPU to one of them 把準備好的process分配給CPU執行
- CPU scheduling decisions may take place when a process: CPU scheduler執行的時間點
  1. Switches from running to waiting state (IO, wait for child)
  2. Switches from running to ready state (timer expire)
  3. Switches from waiting to ready (IO completion)
  4. Terminates

# Non-preemptive vs. 非搶佔 v.s 搶佔

## Preemptive Scheduling

- **Non-preemptive Scheduling/Cooperative Scheduling**
  - Scheduling takes place only under circumstances 1 and 4
  - Process holds the CPU until termination or waiting for IO
  - MS Windows 3.1; Mac OS ( before Mac OS X)
  - Does not require specific HW support for preemptive scheduling
    - E.g., timer
- **Preemptive Scheduling**
  - Scheduling takes place under all the circumstances (1 to 4)
  - Better for time-sharing system and real-time systems
  - Usually, more context switches
  - A cost associated with shared data access
    - May be preempted in an unsafe point

# Scheduling Criteria

- Used to judge the performance of a scheduling algorithm  
好壞的標準
- **CPU utilization**
  - $(100\% - \text{ratio of CPU idle})$
- **Throughput** 吞吐量
  - # of processes that complete their execution per time unit
- **Turnaround time** 周轉時間
  - amount of time to execute a particular process
  - From process submission to process termination

# Scheduling Criteria

- **Waiting time**
  - amount of time a process has been waiting in the ready queue
  - Scheduler does not affect the time for
    - Execution instructions
    - Performing IOs
- Here, we do not consider the memory (including cache) effect*
  - 假設所有記憶題存取時間都一樣
- **Response time**
  - amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)  
    提交request到第一次回應的時間

# Optimization Criteria

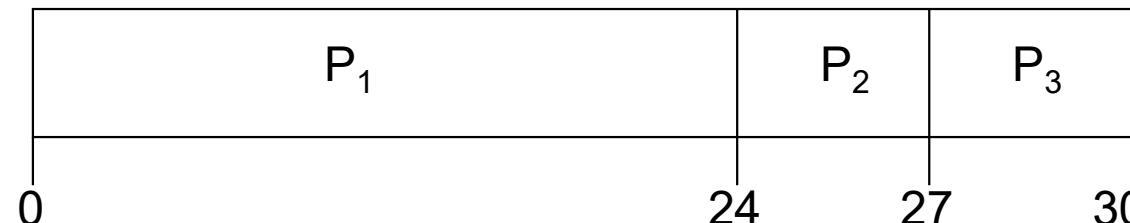
最佳化的目標

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

# First-Come, First-Served (FCFS) Scheduling

Process	Burst Time	← CPU burst
$P_1$	24	
$P_2$	3	
$P_3$	3	

- Implemented via a FIFO queue
- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$   
The Gantt Chart for the schedule is:



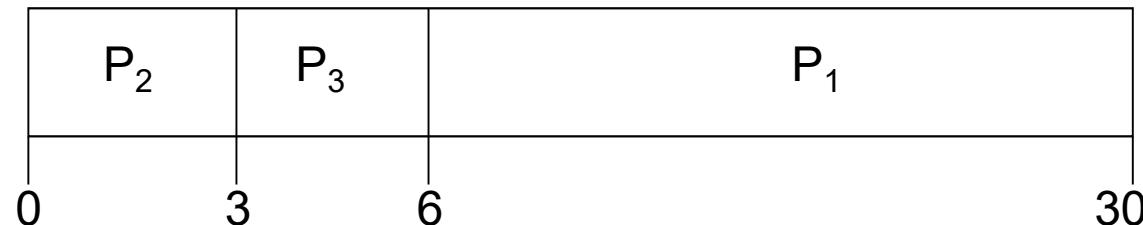
- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$

# FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Much better than previous case

# FCFS Scheduling (Cont.)

- *Convoy effect* 護航效應
  - Short process behind long process
  - Multiple IO bound process may wait for a single CPU bound process  
• Device idle....  
短得排在長的後面，平均等待時間就很長
- FCFS is non-preemptive
  - Not good for time-sharing systems

# Shortest-Job-First (SJF) Scheduling

把process跟下一個CPU burst (下一個

process的開頭) 連接起來，然後選最短的

- Associate with each process the length of its **next** CPU burst, and select the process with the shortest burst to run

- Two schemes

1 – **Non-preemptive** – once the CPU is given to a process, it cannot be preempted until the completion of the CPU burst

2 – **Preemptive** – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt the current process

如果新到達的process的CPU Brust較短，就搶佔，先執行

- known as the **Shortest-Remaining-Time-First (SRTF)** scheduling



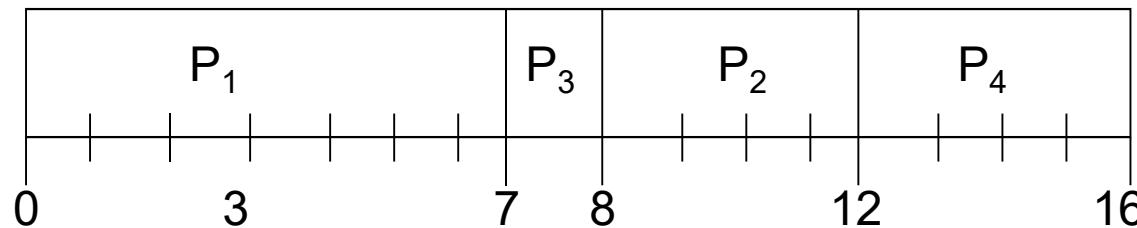
- SJF is optimal – gives minimum average waiting time for a given set of processes

# Example of Non-Preemptive SJF

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- SJF (non-preemptive)

先執行完第一個之後才從後面選  
時間短的先來執行



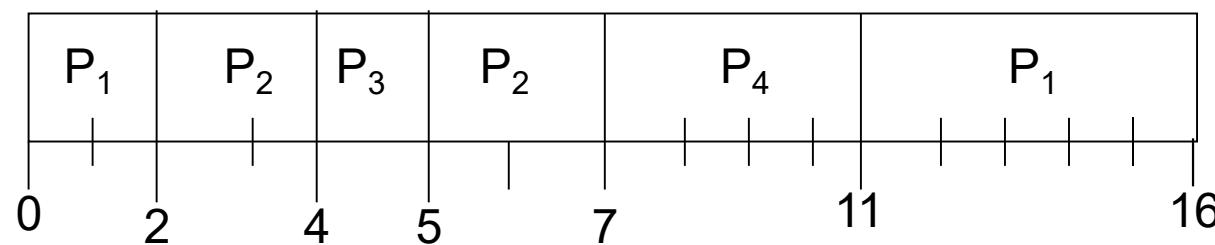
- Average waiting time =  $(0 + 6 + 3 + 7)/4 = 4$

# Example of Preemptive SJF

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- SJF (preemptive)

短的來了之後直接先執行



- Average waiting time =  $(9 + 1 + 0 + 2)/4 = 3$

# SJF Scheduling

- How to know the length of the **next** CPU burst?
  - Difficult.....
  - There is no easy way to know the length of the next CPU burst
  - So, ~~guess~~ predict it ....

# Predicting Length of Next CPU Burst

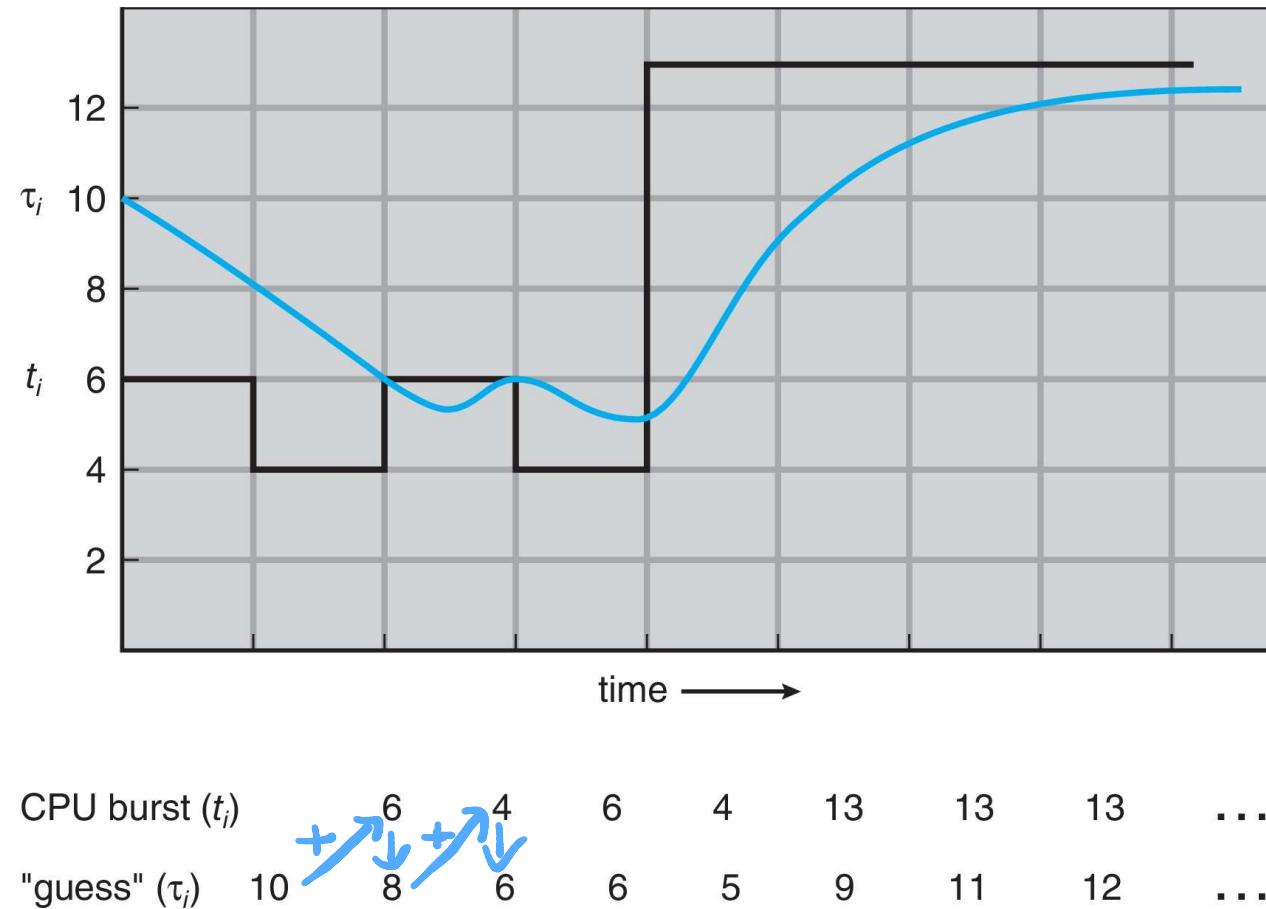
預測下一個CPU Burst的時間

- Can only **estimate** the length
- Can be done by using the length of previous CPU bursts, using exponential averaging
  1.  $t_n$  = actual length of  $n^{th}$  CPU burst
  2.  $\tau_{n+1}$  = predicted value for the next CPU burst
  3.  $\alpha, 0 \leq \alpha \leq 1$
  4. Define:

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

# Predicting Length of the Next CPU Burst

$$\tau_0 = 10$$
$$\alpha = 1/2$$



# Examples of Exponential Averaging

- $\alpha = 0$ 
  - $\tau_{n+1} = \tau_n$
  - Recent history does not count
- $\alpha = 1$ 
  - $\tau_{n+1} = t_n$
  - Only the actual last CPU burst counts
- If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} &= \alpha t_n + (1-\alpha) \alpha t_{n-1} + \dots \\ &\quad + (1-\alpha)^j \alpha t_{n-j} + \dots \\ &\quad + (1-\alpha)^{n+1} \tau_0\end{aligned}$$

- Since both  $\alpha$  and  $(1-\alpha)$  are typically less than 1, each successive term has less weight than its predecessor
  - 展開後可以發現，每一項的權重是逐漸遞減
  - 也就是說  $r_2$  對  $r_3$  的權重影響比較大，相較於  $r_1$  對  $r_3$  來說

# Round Robin (RR)

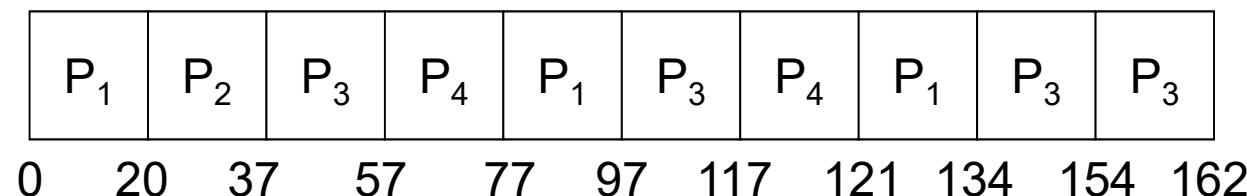
每個process都分到一點CPU

- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.  
可能分到10毫秒或100毫秒
- A process will leave the running state if 執行完之後就會被搶佔，並加回到ready queue
  - Time quantum expire
  - Wait IO or events
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets 1/n of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)q$  time units.
- RR is preemptive

# Example of RR with Time Quantum = 20

Process	Burst Time
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24

- The Gantt chart is:



- • Typically, longer average turnaround time than SJF, but better *response* time

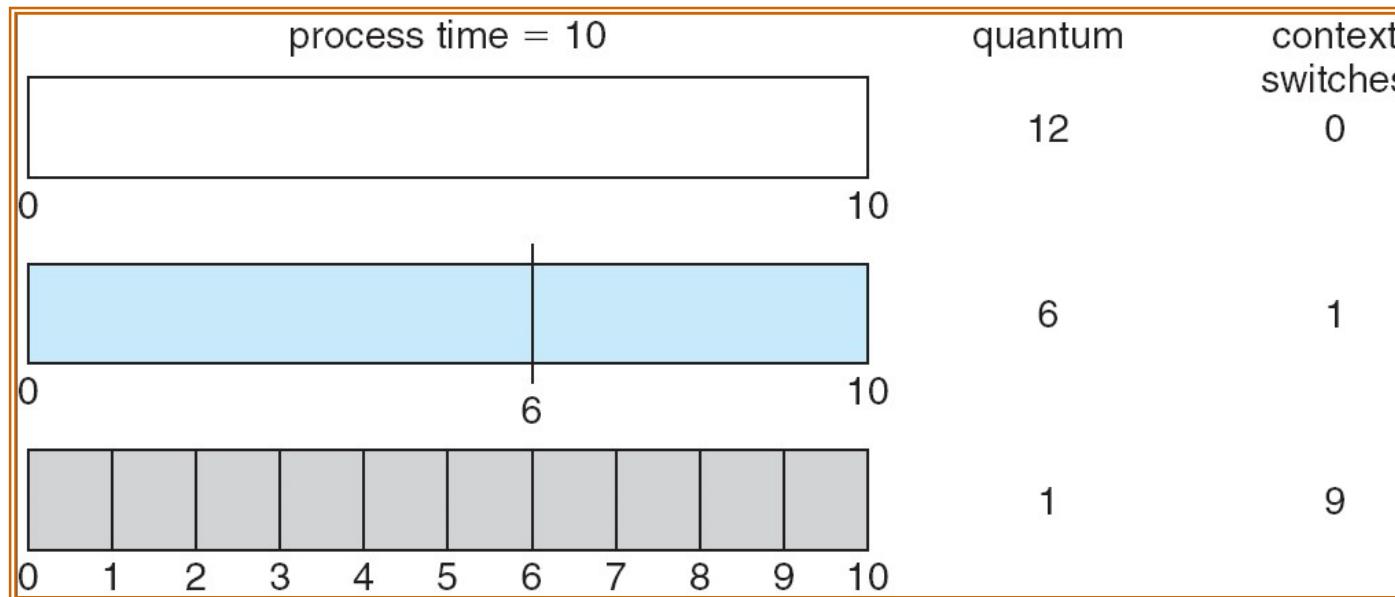
# Time Quantum and Context Switch Counts

## Performance

$q$  large  $\Rightarrow$  FIFO

$q$  small  $\Rightarrow$  a large number of context switches

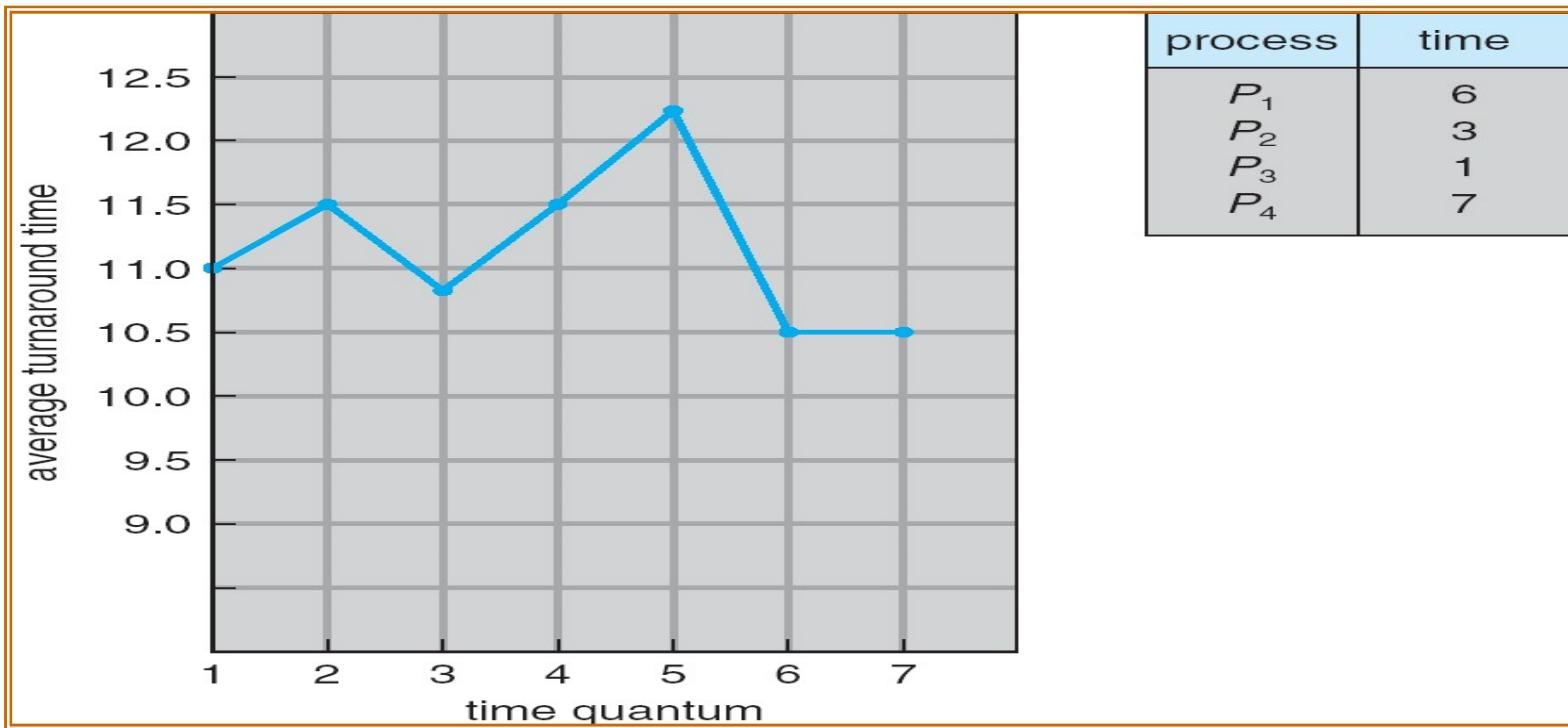
- $q$  must be **large** with respect to context switch time, otherwise overhead is too high



Context switches are not free!!!

Process進入到離開的時間

# Turnaround Time Varies with the Time Quantum



Given 3 processes of 10 time units

for quantum of 1 time unit  $\rightarrow$  average turnaround time = 29

for quantum of 10 time unit  $\rightarrow$  average turnaround time = 20



Rule of thumb: 80% of the CPU bursts should be shorter than the time quantum

經驗總結

# Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the **highest priority** (in many systems, smallest integer → highest priority)
  - Preemptive
  - Non-preemptive

Process	Burst Time	Priority
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2

Execution Sequence: P2, P5, P1, P3, P4

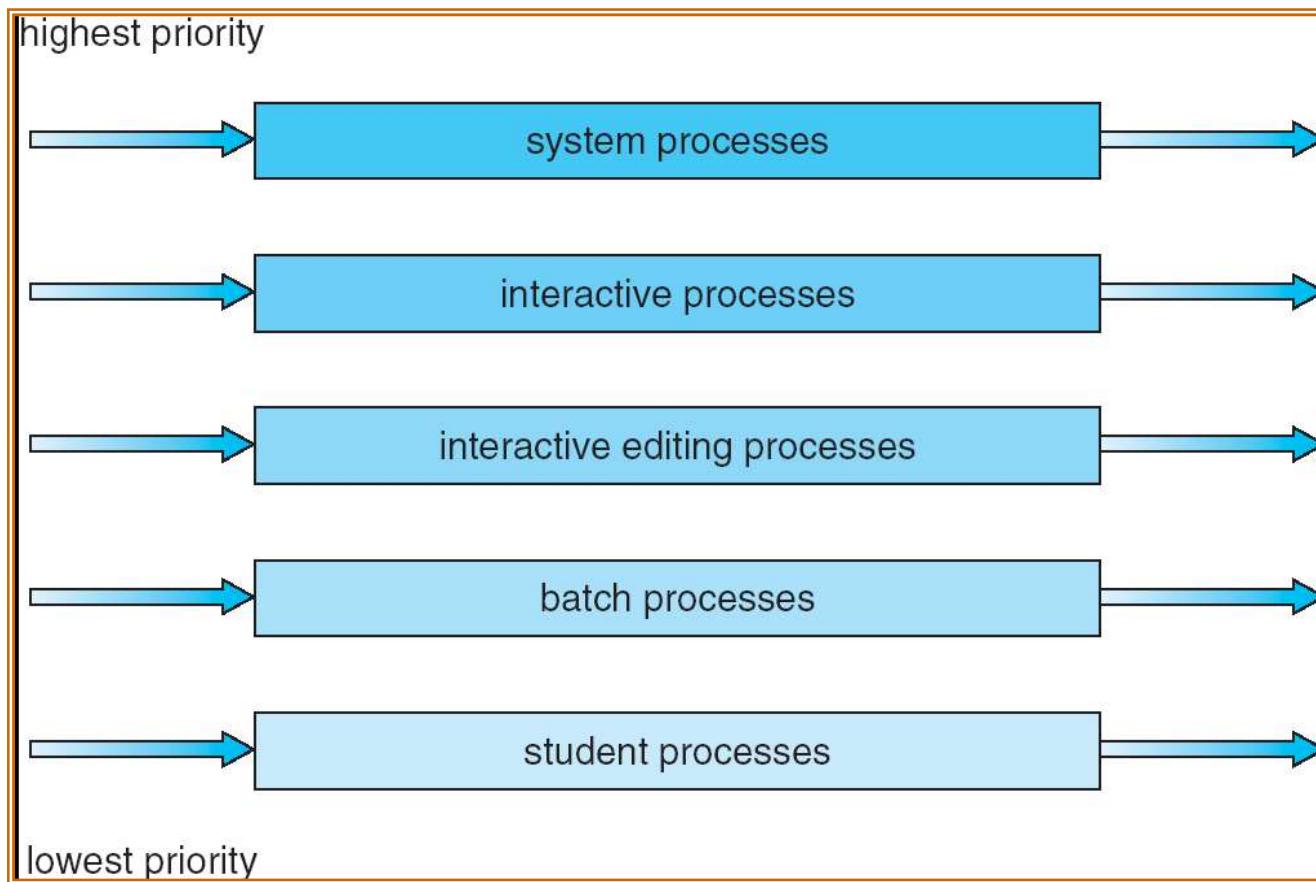
# Priority Scheduling

- The concept is general
  - SJF is a priority scheduling where priority is set according to the predicted next CPU burst time
- Problem : **Starvation** – low priority processes may never execute
  - A low priority process submitted in 1967 had not been run when the system IBM 7094 at MIT was shutdown in 1973
- Solution : Aging – as time progresses increase the priority of the process

# Multilevel Queue

- Used when processes are easily classified into different groups
- Ready queue is partitioned into separate queues
  - e.g., foreground (interactive) and background (batch)
    - These two types of processes have **different response time requirements** 例如分成前端（網頁）和後端（編譯程式）之類的
    - FG processes can have priority over BG processes 前端優先級比後端高
- A process is **fixed** on one queue
- Each queue has its own scheduling algorithm
  - E.g., foreground – RR; background – FCFS

# Multilevel Queue Scheduling



# Multilevel Queue

- Scheduling must be done **between** the queues
  - Fixed priority scheduling 固定優先級
    - i.e., serve all from foreground then from background
    - possibility of starvation 可能有人都不被執行到
  - Time slice 時間切分
    - each queue gets a certain amount of CPU time which it can schedule amongst its processes
    - i.e., 80% to foreground in RR, 20% to background in FCFS 消除starvation

# Multilevel Feedback Queue

- A process can move among different queues
- The idea 這個process會穿梭在queues之間
  - Separate processes according to the characteristics of their **CPU bursts** 佔用太多CPU 就會被分類到低優先
    - Use too much CPU time → move to a lower priority Q
      - Favor interactive and IO bound processes

解釋：

- 核心機制： 這是 MLFQ 與靜態多級佇列最大的區別。行程的優先級（即它所在的佇列）是動態的。
- 回饋 (Feedback)： 「回饋」一詞的含義是：系統會根據行程在執行過程中的實際表現（即回饋），來動態調整其優先級。



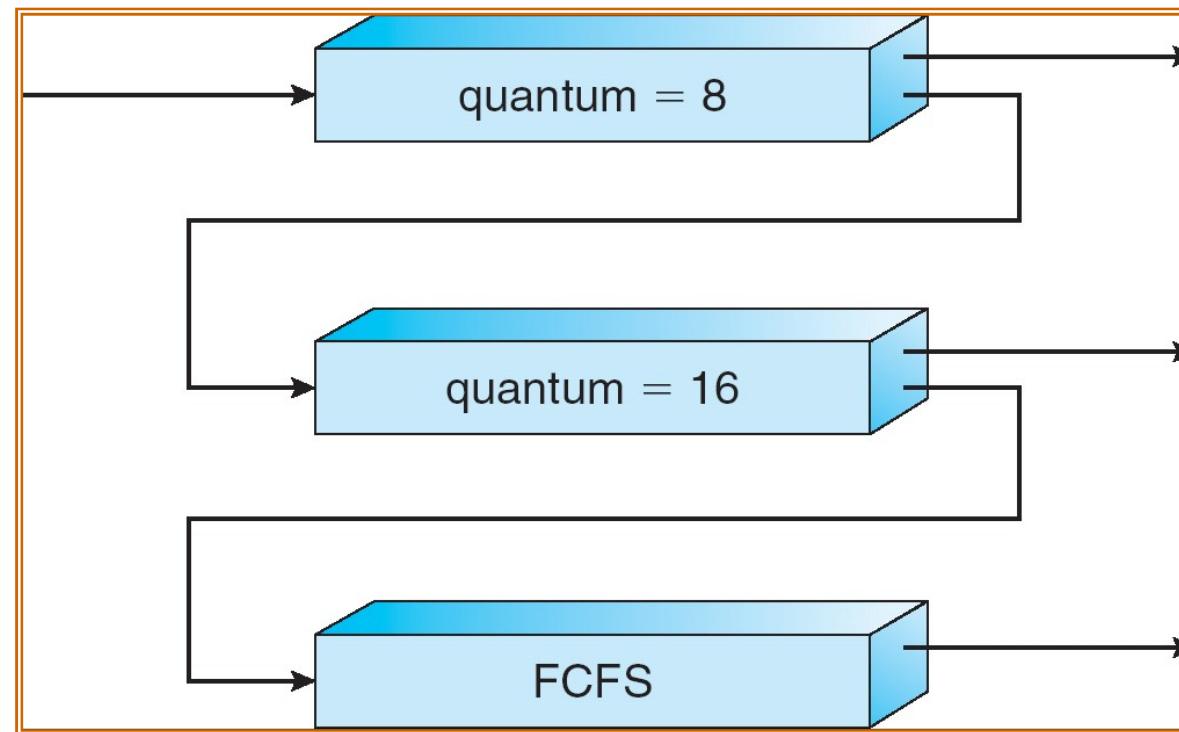
解釋：

- 目標： MLFQ 的主要目標是自動識別出行程是 **CPU 密集型**（長 CPU 突發）還是 **I/O 密集型**（短 CPU 突發）。
- 動態分類： MLFQ 不需要像 SJF 那樣去「預測」CPU 突發，而是透過行程的實際執行情況來事後分類。

# Example of Multilevel Feedback Queue

- Three queues:
  - $Q_0$  – RR with time quantum 8 milliseconds
  - $Q_1$  – RR time quantum 16 milliseconds
  - $Q_2$  – FCFS
- Scheduling
  - A new job enters queue  $Q_0$ . When it gains CPU, job receives 8 milliseconds. If it does not finish its **current burst** in 8 milliseconds, job is **preempted** and moved to queue  $Q_1$ .  
搶佔
  - At  $Q_1$  job is again served and receives 16 additional milliseconds. If it still does not complete **its burst**, it is **preempted** and moved to queue  $Q_2$ .

# Multilevel Feedback Queues



**Give highest priority to processes with CPU burst  $\leq 8\text{ms}$**

# Multilevel Feedback Queues

- Multilevel-feedback-queue scheduler defined by the following parameters:
  - 1 – number of queues 序列數量 每個序列不同的演算法 (FCFS,RR之類)
  - 2 – scheduling algorithms for each queue 提升process優先級的方式
  - 3 – method used to determine when to upgrade a process 降低process優先級的方式
  - 4 – method used to determine when to demote a process
  - 5 – method used to determine which queue a process will enter when that process needs service 一個process創建後要進入哪個queue
- It is the most generic algorithm
  - Can be configured to match a specific system
- It is the most complex algorithm
  - You have to select a proper value for each parameter

# Multiple-Processor Scheduling

多處理器排程

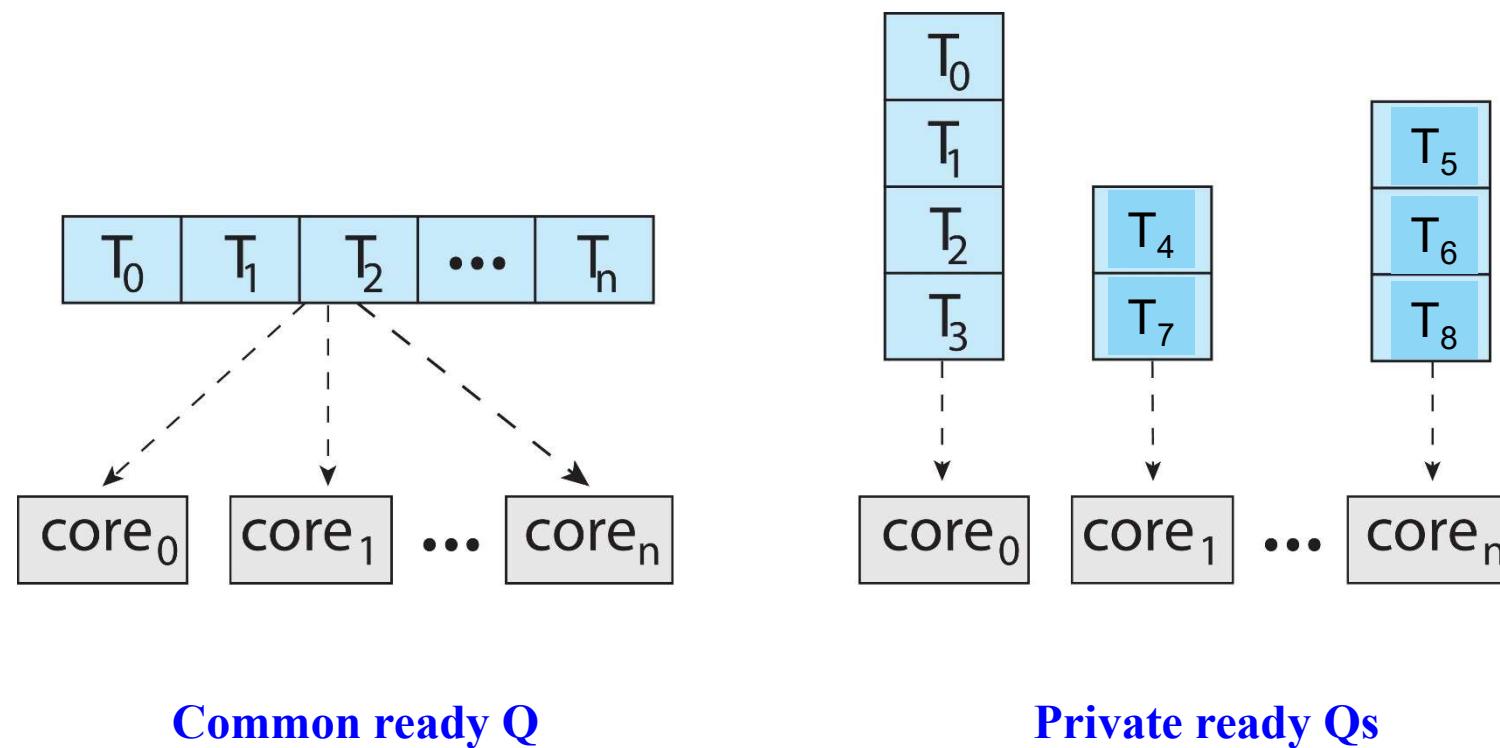
- Load sharing 負載共享
- CPU scheduling is more complex when multiple CPUs are available  
都是同質的處理器
- We consider **homogeneous processors**
  - Can use any **available** processor to run any ready processes
- Topics
  - ASMP vs. SMP
  - Processor affinity
  - Load balancing
  - Multithreaded core

# ASMP vs. SMP

非對稱 vs. 對稱

- Approaches to MP scheduling
  - **Asymmetric multiprocessing (ASMP)**
    - Only one processor accesses the OS data structures, alleviating the need for data sharing 指定其中一個CPU為主控，只有他可以access OS
    - The other processors run user code only
  - **Symmetric multiprocessing (SMP)**
    - All processors can access the OS data structures
    - Each processor is self scheduling 所有CPU都可以操控OS
    - Common or private ready queue (*see next slide*)
    - Scheduler in each processor selects a process from the ready Q
      - In case of common ready Q, must ensure 也就是每個CPU自己的scheduler
        - » Two processors don't choose the same process
        - » Processes are not lost from the Q 去ready queue拿process
    - All modern OSs supports SMP
      - Windows, Linux, Solaris, Mac OS X...
  - *We focus on SMP systems here*

# Common/Private Ready Queues



Common ready Q

Private ready Qs

# Processor Affinity

親和性

- Cache miss rate increases if a process migrates to another processor  
CPU1換到CPU2會造成cache miss  
CPU1的cache可能在CPU2裡面不適用
- Most SMP systems try to avoid migration
  - Processor affinity
    - Keep the process running on the same processor
- Soft affinity
  - Try to keep the process always on a fixed processor
  - But, NO guarantee...
- Hard affinity
  - Guarantee to keep a process always on a fixed processor
  - Linux provides system calls to support hard affinity

# Load Balancing

- Balance the load among the processors
- Only necessary on private-ready-Q systems
  - Different ready Qs can have different lengths
  - In common-ready-Q systems, the load is already balanced
  - Most contemporary OSs use private ready Qs
- Two general approaches
  - | – Push migration
    - a specific task periodically checks the load and balance the load if it finds an imbalance
  - 2 – Pull migration
    - An idle processor pulls a ready task from a busy processor
- The above two approaches can co-exist
  - Linux supports both (Note: It performs push migration every 200ms)

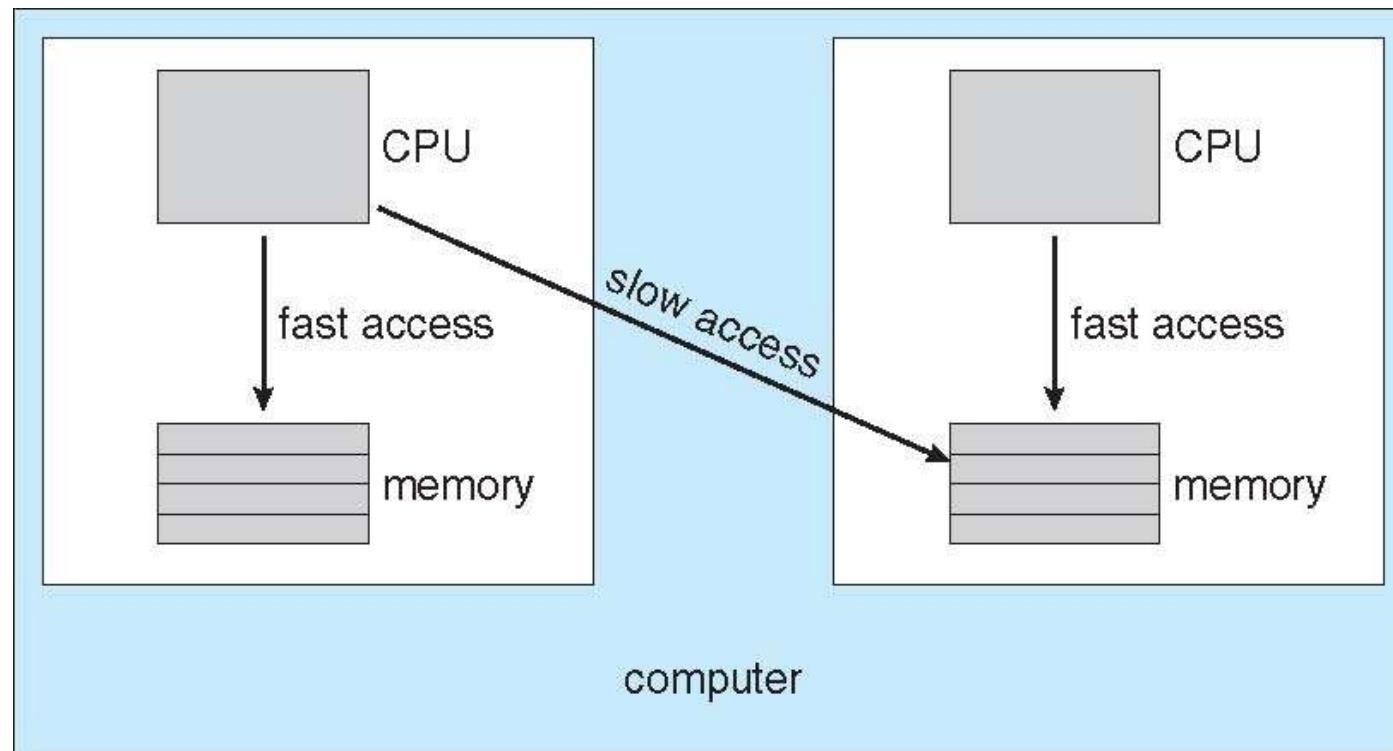
# Load Balancing

- Load balancing often counteract the benefits of processor affinity
  - Load balancing is done by process migration
  - Processor affinity try not to migrate processes
  - An idle processor can
    - Pull processes only when imbalance exceeds a certain threshold

相違背

# NUMA and CPU Scheduling

NUMA architecture also has affinity and load balancing issues...



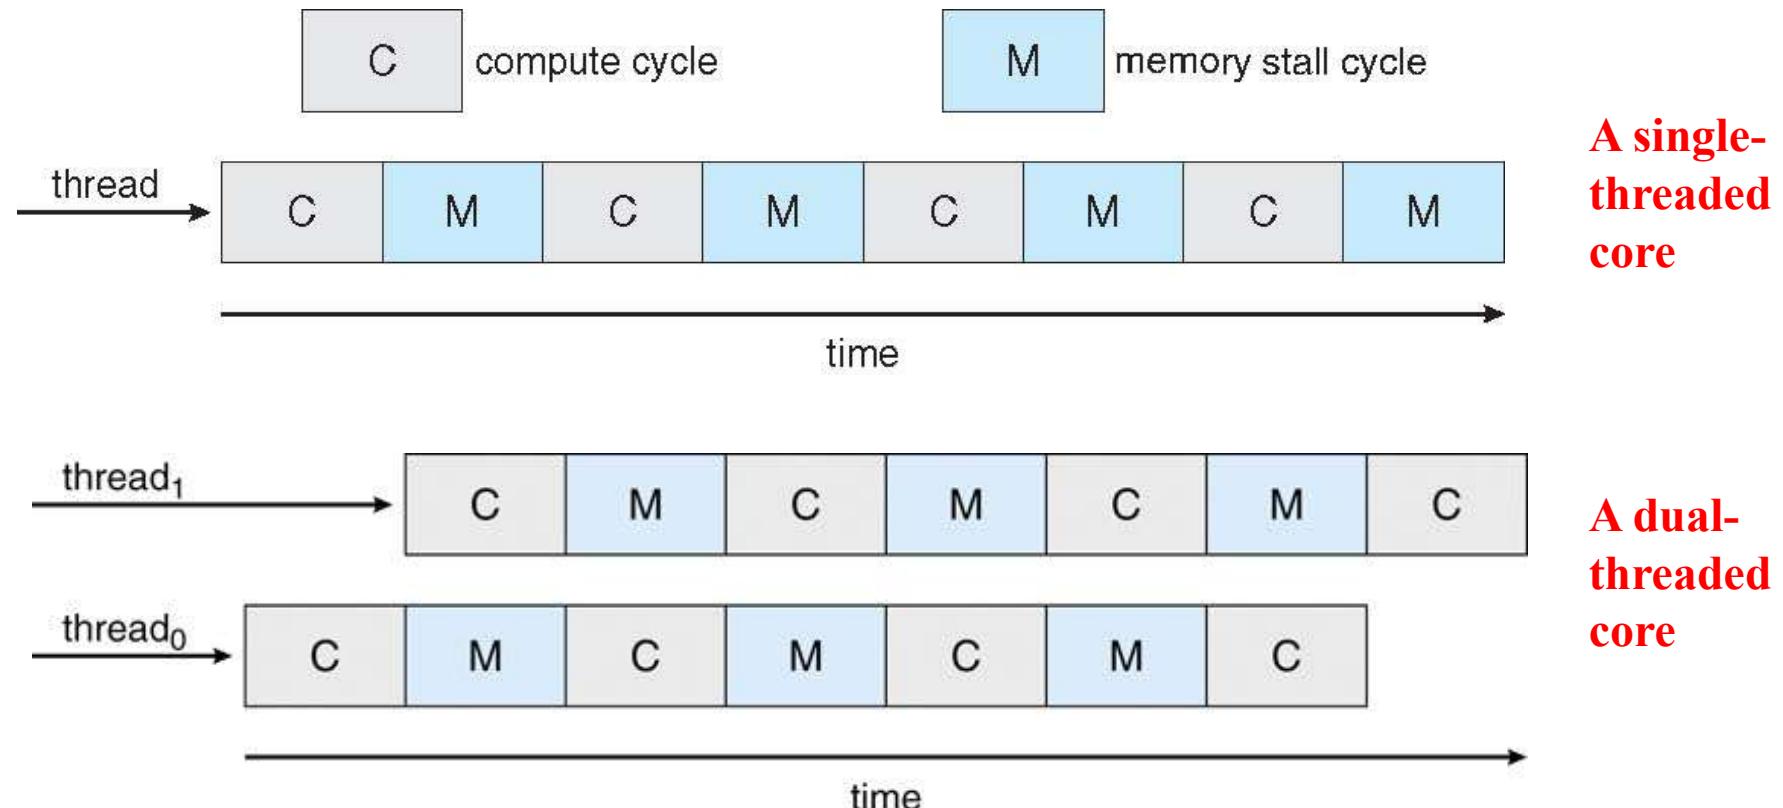
NUMA也有load balancing 的問題

# Multicore Processors

- Recent trend to place multiple processor cores on same physical chip
  - Faster and consume less power
- Multiple (hardware) threads per core also growing
  - Provide multiple logical (not physical) processors on the same physical core (see next slide)
    - Each logical P has its own architecture state
      - General and status registers
    - Each logical P handle its own interrupts
    - Logical Ps share the resources of the physical P such as ALU, cache, FPU..
    - E.g. Intel's hyperthreading technology
  - **Example:** takes advantage of memory stall to make progress on another thread while memory retrieve happens

同一個晶片，分成不同的概念上的processors

# A Multithreading Example in a Core



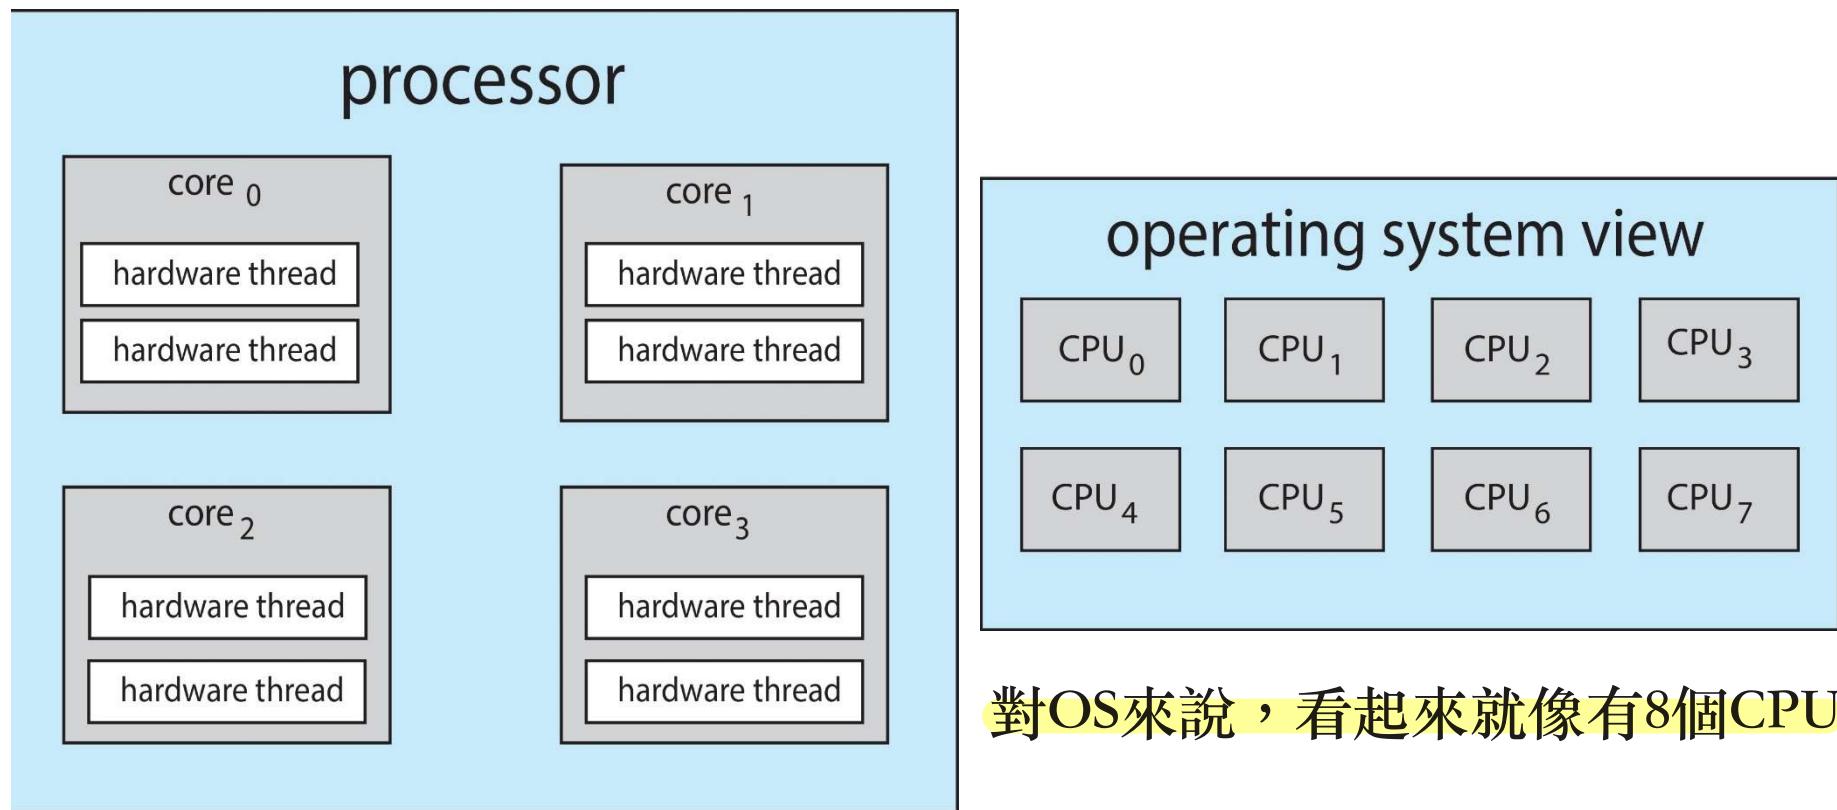
**Temporal** multithreading:

[https://en.wikipedia.org/wiki/Temporal\\_multithreading](https://en.wikipedia.org/wiki/Temporal_multithreading)

**Simultaneous** multithreading (SMT):

[https://en.wikipedia.org/wiki/Simultaneous\\_multithreading](https://en.wikipedia.org/wiki/Simultaneous_multithreading)

# A Multithreaded Multicore System



每個core跑兩個thread所以 $4 \times 2 = 8$

# Real-Time Scheduling

- *Hard real-time* systems – required to complete a critical task within a **guaranteed** amount of time  
在一定時間內完成關鍵任務
- *Soft real-time* computing – requires that critical processes receive priority over the others; **NO guarantee** on the execution time limit  
讓關鍵任務有高優先級

# Thread Scheduling

- 1 • Kernel threads are scheduled by OS
  - 2 • User threads are managed by thread library
- 
- 1 • Local Scheduling – How the threads library decides which thread to put onto an available LWP (kernel thread)
    - For M:1 or M:M models

User  
要在哪個user thread放上  
LWP(kernel 跟 user的橋樑)
  - 2 • Global Scheduling – How the kernel decides which kernel thread to run next  
Kernel

# Contention Scope

競爭CPU

- Process Contention Scope (PCS)
  - Competitions among threads of the same process
  - Scheduling is typically done according to priority
    - Thread priorities are set by programmers, not adjusted by thread lib 多個user競爭較少的kernel  $M:M$
    - Usually no time slicing among threads of equal priority
- System Contention Scope (SCS) 是1:1的時候就只有SCS
  - Competitions among threads in the system
  - Systems with 1:1 model only use SCS
    - Linux, MacOS...

1個 user → 1個 kernel

# Pthread Scheduling

POSIX

- Contention Scope
  - *SCS* – PTHREAD\_SCOPE\_SYSTEM
  - *PCS* – PTHREAD\_SCOPE\_PROCESS
- On M:M systems
  - PTHREAD\_SCOPE\_PROCESS schedules the user thread onto available (and shared) LWPs
    - # of LWPs is determined by the thread lib
  - PTHREAD\_SCOPE\_SYSTEM will bind the user thread to a dedicated LWP
    - Becomes 1:1
- API
  - pthread\_attr\_setscope()
  - pthread\_attr\_getscope()

# Pthread Scheduling API

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5
int main(int argc, char *argv[])
{
    int i;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* set the scheduling algorithm to PROCESS or SYSTEM */
    pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);
    /* set the scheduling policy - FIFO, RR, or OTHER */
    pthread_attr_setschedpolicy(&attr, SCHED_OTHER);
    /* create the threads */
    for (i = 0; i < NUM_THREADS; i++)
        pthread_create(&tid[i], &attr, runner, NULL);
```

# Pthread Scheduling API

```
/* now join on each thread */
for (i = 0; i < NUM_THREADS; i++)
    pthread_join(tid[i], NULL);
} /* end of main() */

/* Each thread will begin control in this function */
void *runner(void *param)
{
    printf("I am a thread\n");
    pthread_exit(0);
}
```

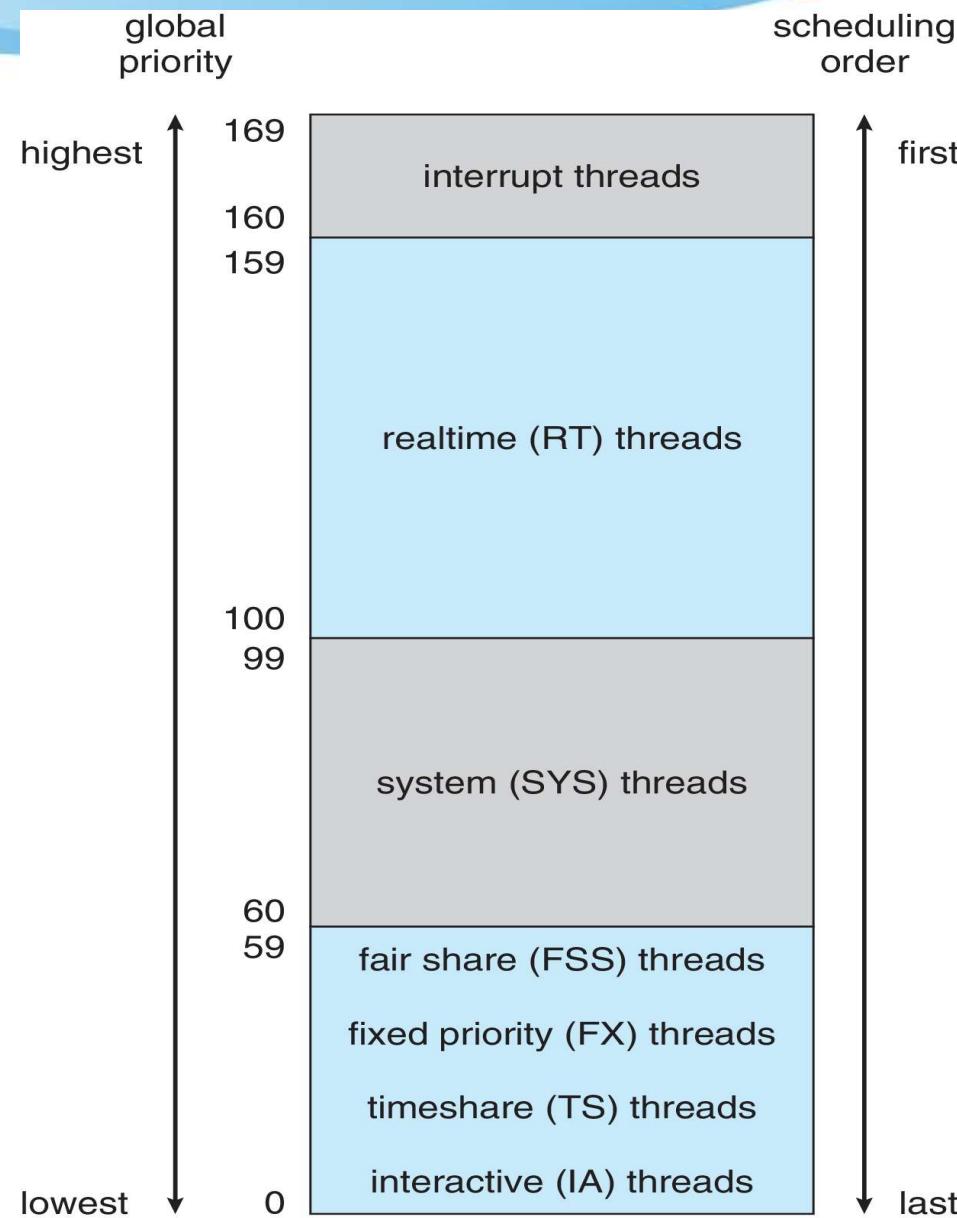
# Operating System Examples

- We mention **kernel thread scheduling** here
- Solaris scheduling
- Windows XP scheduling
- Linux scheduling

# Solaris Scheduling

- Priority based
  - RR for same-priority threads
- 6 classes
  - Real time (RT)
  - System (SYS)
  - Fair Share (FSS)
  - Fixed priority (FP)
  - Time sharing (TS) -- default
  - Interactive (IA)

# Solaris Scheduling



# Solaris Scheduling

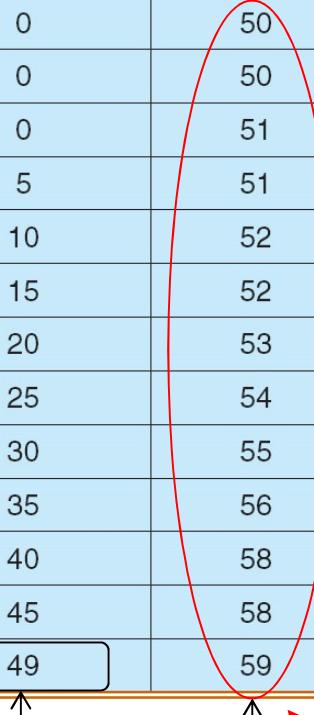
- Real time (RT) class
  - The highest priority among the 6 classes
  - Allows a RT process to have fast responses
- System class
  - Kernel processes, such as paging daemon
- TS/IA classes
  - Dynamically alters priorities
  - Assign time slices of different lengths using a **multilevel feedback Q**
    - Higher priority → smaller time slice
      - Good response time for interactive processes
      - Good throughput for CPU-bound processes

# Solaris Dispatch Table for TS/IA Classes

priority	time quantum	time quantum expired	return from sleep
0	200	0	50
5	200	0	50
10	160	0	51
15	160	5	51
20	120	10	52
25	120	15	52
30	80	20	53
35	80	25	54
40	40	30	55
45	40	35	56
50	40	40	58
55	40	45	58
59	20	49	59

Lowest priority →

Priority change to favor IO bound processes



# Windows XP Scheduling

priority-based preemptive scheduling

	real-time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1

Priority  
classes

Relative  
Priority in  
a class

Increase the quantum of the **foreground** process by some factor (e.g., 3)

# Linux Scheduling

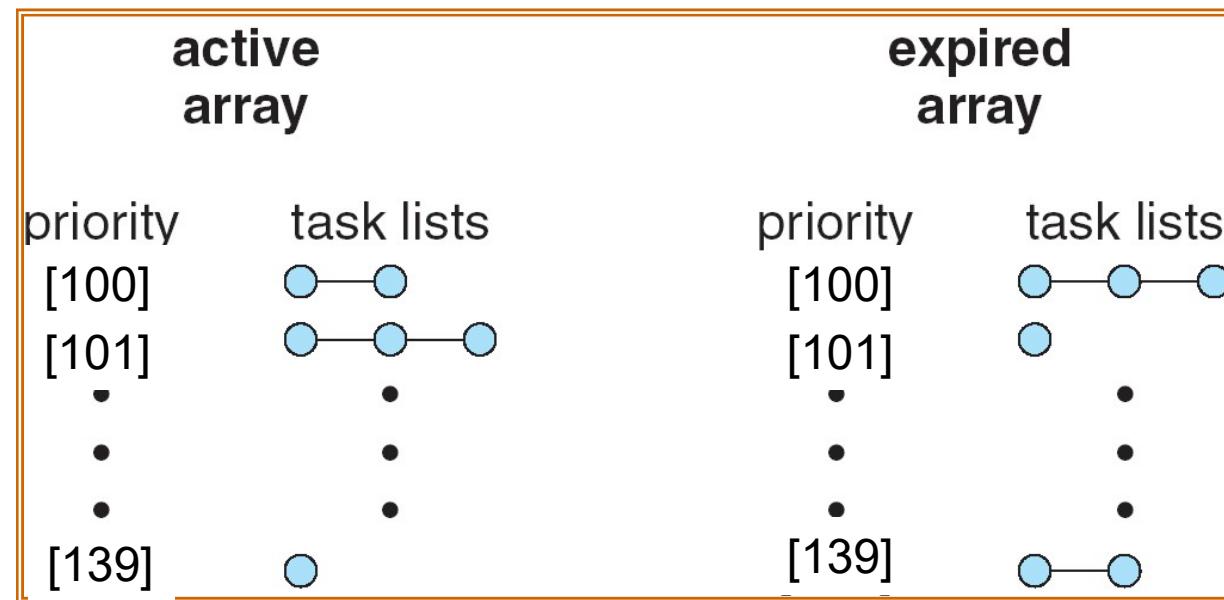
- Two algorithms: time-sharing and real-time (**soft**)
- Time-sharing
  - O(1) scheduler (kernel 2.5)
    - Prioritized & credit-based
      - Credit subtracted when timer interrupt occurs
      - When credit = 0, another process chosen
      - When all processes have credit = 0, re-crediting occurs
        - » Based on factors including priority and history
    - Priority boosts for interactive or IO bound processes
  - CFS (after kernel 2.6)
- Real-time
  - Soft real-time
  - POSIX.1b compliant (IEEE 1003.1b-1993) – two classes
    - FCFS and RR
    - Highest priority process always runs first

# O(1) Scheduler - Relationship between Priorities and Time-slice Length

numeric priority	relative priority	time quantum
0	highest	200 ms
•		
•		
•		
99		
100		
•		
•		
•		
[139]	lowest	10 ms

# O(1) Scheduler - List of Tasks Indexed According to Priorities

Each ready Q contains two arrays



# Linux CFS Scheduler (after kernel 2.6)

- Completely Fair Scheduler (CFS)
- Schedule task with the smallest score
  - Derived from **virtual run time** of the task
  - Tasks with the smallest virtual run time tend to be selected to run
- **nice** value can affect the score
  - $\text{nice} > 0$  (lower priority) → increase score
  - $\text{nice} < 0$  (higher priority) → decrease score

期中考到這

# Java Thread Scheduling

- JVM uses a **Preemptive, Priority-based** scheduling algorithm
- **FIFO** queue is used if there are multiple threads with the **same priority**

# Java Thread Scheduling (cont.)

JVM Schedules a Thread to Run When:

1. The currently running thread exits the Runnable state
2. A higher priority thread enters the Runnable state

JVM doesn't ensure time-slicing

# Time-Slicing

Since the JVM doesn't ensure time-slicing, the `yield()` method may be used:

```
while (true) {  
    // perform CPU-intensive task  
    ...  
    Thread.yield();  
}
```

This yields control to another thread of **equal priority**

# Thread Priorities

<u>Priority</u>	<u>Comment</u>
Thread.MIN_PRIORITY	Minimum Thread Priority
Thread.MAX_PRIORITY	Maximum Thread Priority
Thread.NORM_PRIORITY	Default Thread Priority

Priorities may be set by using the **setPriority()** method:

```
setPriority(Thread.NORM_PRIORITY + 2);
```

# Algorithm Evaluation

- Deterministic modeling
  - takes a particular **predetermined workload** and defines the performance of each algorithm for that workload
  - Simple and fast
    - Useful only when the set of programs and their behaviors are fixed
- Queueing models
  - Assumes the **distribution** of the **burst length** and **process arrival rates**
  - It's possible to compute the average throughput, utilization, waiting time...

# Algorithm Evaluation

- Simulation
  - *see next slide*
  - Still of limited accuracy
- Implementation
  - The only completely accurate way to evaluate an algorithm
  - High cost

# Simulation

